

Real-Estate Professionals Scraper + ML + Deep Learning

1. Overview

This project collects **LinkedIn profiles of real estate professionals in India** using Selenium-based scraping, then applies **two machine-learning pipelines** to automatically detect whether a profile is **relevant to the real estate domain**:

1. **Classical ML Baseline** – TF-IDF vectoriser + Logistic Regression.
2. **Deep-Learning Model** – Keras **Embedding + Bi-LSTM** network.

The pipeline produces:

- **Raw scraped CSV** of all profiles.
- **Scored CSV** with ML & DL probabilities and predicted relevance labels.
- **Filtered CSV** with only high-confidence “relevant” profiles.
- **Uncertain CSV** with mid-confidence profiles for manual review.
- A **ZIP archive** for easy sharing.
- Saved **vectoriser / ML model / DL model/tokenizer** for later reuse.

The code is designed to **recover gracefully from LinkedIn session drops** and to **handle very small datasets**.

2. Usage

1. Configure credentials

Edit these two variables near the top of the script:

2. `LINKEDIN_EMAIL = "your_email@example.com"`
3. `LINKEDIN_PASSWORD = "your_password"`

4. Run the script

5. `python real_estate_scraper_ml_dl.py`

6. Outputs (all saved in the working directory):

- `real_estate_professionals_india_scored.csv` – all profiles with ML + DL scores
- `real_estate_professionals_india_filtered.csv` – only high-probability relevant profiles
- `real_estate_professionals_india_uncertain.csv` – medium-confidence profiles
- `real_estate_professionals_india.csv` – identical to *filtered.csv* (legacy name)
- `real_estate_professionals_india.zip` – zipped filtered CSV
- `models/` folder with:
 - `relevance_vectorizer.joblib`
 - `relevance_model.joblib`
 - `dl_tokenizer.joblib`
 - `dl_model.h5`

7. Typical workflow

- The script first generates random (company, city, job title) queries.
 - Launches Chrome, logs into LinkedIn.
 - Scrapes profiles per query, deduplicates them.
 - Builds weak labels using keyword-based rules.
 - **Trains or loads** both the baseline and DL models.
 - Scores every profile with **probabilities** and **predicted labels**.
 - Saves final CSVs and ZIP for distribution.
-

3. Rationale for Model Choices

3.1 Weak-Label + ML Approach

- Scraped data is **unlabelled**.
- A **rule-based weak-labelling** step (keywords such as “broker”, “property manager”) provides **noisy binary labels** to bootstrap training.

3.2 TF-IDF + Logistic Regression (Baseline)

- **Interpretable & light-weight**: good for small datasets and explainable terms.
- **TF-IDF** captures important words & bigrams in titles/positions.
- **Logistic Regression** is **fast to train**, robust with sparse data, and supports **class-imbalance weighting**.

3.3 Deep-Learning Bi-LSTM

- For larger datasets, **word order and context** can be informative.
- The **Embedding + Bidirectional LSTM** network captures **sequential patterns** that TF-IDF ignores.
- Uses **Early-Stopping** and neutral fall-backs to avoid over-fitting or failures on small data.

3.4 Hybrid Output

- Both models’ **probabilities** are retained:

prob_relevant – baseline TF-IDF + LogReg probability.

dl_prob_relevant – Deep-Learning probability.

- Users can compare or ensemble them later.
-

4. Robustness & Safety Features

- **Validation-split auto-adjustment:** skips validation if <10 samples.
 - **Dummy-model fall-back:** if dataset <2 samples, DL model is skipped and outputs a neutral probability (0.5) for each profile.
 - **Shape-mismatch guard:** fixes the Keras output-length bug encountered with tiny datasets.
 - **Auto-restart of the Selenium driver** if the LinkedIn session expires.
 - **Deduplication** to avoid repeated profiles.
-